

```
### Name: Naheed Sajjad
### Student ID:A00326697
### FinalProject_MachineLearning
```

```
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import numpy as np # linear algebra
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
import copy
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix,
## Data used "heart_2022_with_nans.csv" downloaded from "https://www.kaggle.com/datasets/kamilpytlak/personal-key-ind
```

```
## Step 1: Load the dataset
df = pd.read_csv("heart_2022_with_nans.csv")
# Display first few rows
df.head()
```

	State	Sex	GeneralHealth	PhysicalHealthDays	MentalHealthDays	LastCheckupTime	PhysicalActivities	SleepHours
0	Alabama	Female	Very good	0.0	0.0	Within past year (anytime less than 12 months ...)	No	8.0
1	Alabama	Female	Excellent	0.0	0.0	NaN	No	6.0
2	Alabama	Female	Very good	2.0	3.0	Within past year (anytime less than 12 months ...)	Yes	5.0
3	Alabama	Female	Excellent	0.0	0.0	Within past year (anytime less than 12 months ...)	Yes	7.0
4	Alabama	Female	Fair	2.0	0.0	Within past year (anytime less than 12 months ...)	Yes	9.0

5 rows × 40 columns

```
## Step 2: Understand Dataset Structure
# Check dataset shape (no of rows*no of columns)
df.shape
(445132, 40)
```

```
# Display column names
df.columns
Index(['State', 'Sex', 'GeneralHealth', 'PhysicalHealthDays',
       'MentalHealthDays', 'LastCheckupTime', 'PhysicalActivities',
       'SleepHours', 'RemovedTeeth', 'HadHeartAttack', 'HadAngina',
       'HadStroke', 'HadAsthma', 'HadSkinCancer', 'HadCOPD',
       'HadDepressiveDisorder', 'HadKidneyDisease', 'HadArthritis',
       'HadDiabetes', 'DeafOrHardOfHearing', 'BlindOrVisionDifficulty',
       'DifficultyConcentrating', 'DifficultyWalking',
       'DifficultyDressingBathing', 'DifficultyErrands', 'SmokerStatus',
       'ECigaretteUsage', 'ChestScan', 'RaceEthnicityCategory', 'AgeCategory',
       'HeightInMeters', 'WeightInKilograms', 'BMI', 'AlcoholDrinkers',
       'HIVTesting', 'FluVaxLast12', 'PneumoVaxEver', 'TetanusLast10Tdap',
       'HighRiskLastYear', 'CovidPos'],
      dtype='object')
```

```
## Step 3: Handle Missing Values
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 445132 entries, 0 to 445131
Data columns (total 40 columns):
 #   Column            Non-Null Count  Dtype 

```

```
0 State 445132 non-null object
1 Sex 445132 non-null object
2 GeneralHealth 443934 non-null object
3 PhysicalHealthDays 434205 non-null float64
4 MentalHealthDays 436065 non-null float64
5 LastCheckupTime 436824 non-null object
6 PhysicalActivities 444039 non-null object
7 SleepHours 439679 non-null float64
8 RemovedTeeth 433772 non-null object
9 HadHeartAttack 442067 non-null object
10 HadAngina 440727 non-null object
11 HadStroke 443575 non-null object
12 HadAsthma 443359 non-null object
13 HadSkinCancer 441989 non-null object
14 HadCOPD 442913 non-null object
15 HadDepressiveDisorder 442320 non-null object
16 HadKidneyDisease 443206 non-null object
17 HadArthritis 442499 non-null object
18 HadDiabetes 444045 non-null object
19 DeafOrHardOfHearing 424485 non-null object
20 BlindOrVisionDifficulty 423568 non-null object
21 DifficultyConcentrating 420892 non-null object
22 DifficultyWalking 421120 non-null object
23 DifficultyDressingBathing 421217 non-null object
24 DifficultyErrands 419476 non-null object
25 SmokerStatus 409670 non-null object
26 ECigaretteUsage 409472 non-null object
27 ChestScan 389086 non-null object
28 RaceEthnicityCategory 431075 non-null object
29 AgeCategory 436053 non-null object
30 HeightInMeters 416480 non-null float64
31 WeightInKilograms 403054 non-null float64
32 BMI 396326 non-null float64
33 AlcoholDrinkers 398558 non-null object
34 HIVTesting 379005 non-null object
35 FluVaxLast12 398011 non-null object
36 PneumoVaxEver 368092 non-null object
37 TetanusLast10Tdap 362616 non-null object
38 HighRiskLastYear 394509 non-null object
39 CovidPos 394368 non-null object
```

dtypes: float64(6), object(34)
memory usage: 135.8+ MB

```
# Check for missing values
df.isnull().sum() #No missing value detected.
```

	0
State	0
Sex	0
GeneralHealth	1198
PhysicalHealthDays	10927
MentalHealthDays	9067
LastCheckupTime	8308
PhysicalActivities	1093
SleepHours	5453
RemovedTeeth	11360
HadHeartAttack	3065
HadAngina	4405
HadStroke	1557
HadAsthma	1773
HadSkinCancer	3143
HadCOPD	2219
HadDepressiveDisorder	2812
HadKidneyDisease	1926
HadArthritis	2633
HadDiabetes	1087
DeafOrHardOfHearing	20647
BlindOrVisionDifficulty	21564
DifficultyConcentrating	24240
DifficultyWalking	24012
DifficultyDressingBathing	23915
DifficultyErrands	25656
SmokerStatus	35462
ECigaretteUsage	35660
ChestScan	56046
RaceEthnicityCategory	14057
AgeCategory	9079
HeightInMeters	28652
WeightInKilograms	42078
BMI	48806
AlcoholDrinkers	46574
HIVTesting	66127
FluVaxLast12	47121
PneumoVaxEver	77040
TetanusLast10Tdap	82516
HighRiskLastYear	50623
CovidPos	50764

dtype: int64

```
# remove any rows that contain missing values (NaN) and to modify the DataFrame directly.
df.dropna(inplace=True)
```

```
## Step 4: Check for Duplicate Records
# Check for duplicates
```

```
df.duplicated().sum()
```

```
np.int64(9)
```

```
#parameters keep the first occurrence of each duplicated row and remove subsequent.  
# Performs the operation directly on the existing DataFrame.  
df.drop_duplicates(keep='first', inplace=True)
```

```
df.shape
```

```
(246013, 40)
```

```
## Step 5: Summary Statistics of Data  
df.describe()
```

	PhysicalHealthDays	MentalHealthDays	SleepHours	HeightInMeters	WeightInKilograms	BMI
count	246013.000000	246013.000000	246013.000000	246013.000000	246013.000000	246013.000000
mean	4.119055	4.167292	7.021312	1.705150	83.615522	28.668258
std	8.405803	8.102796	1.440698	0.106654	21.323232	6.514005
min	0.000000	0.000000	1.000000	0.910000	28.120000	12.020000
25%	0.000000	0.000000	6.000000	1.630000	68.040000	24.270000
50%	0.000000	0.000000	7.000000	1.700000	81.650000	27.460000
75%	3.000000	4.000000	8.000000	1.780000	95.250000	31.890000
max	30.000000	30.000000	24.000000	2.410000	292.570000	97.650000

```
#Check for unique values in each column  
for col in df.columns:  
    print(f"Value counts for column '{col}':")  
    print(df[col].value_counts())  
    print("-" * 30)      #repeats the string "-" 30 times.
```

```

res, received tetanus shot, but not trap
Name: count, dtype: int64
-----
Value counts for column 'HighRiskLastYear':
HighRiskLastYear
No      235437
Yes     10576
Name: count, dtype: int64
-----
Value counts for column 'CovidPos':
CovidPos
No          167297
Yes         70324
Tested positive using home test without a health professional    8392
Name: count, dtype: int64
-----
```

```

## Step-6: Exploratory Data Analysis Visualizations
#df["HadHeartAttackNum"] = df["HadHeartAttack"].map({"No": 0, "Yes": 1})
#df['HadHeartAttack'].value_counts().plot(kind = 'bar')
heart_col = "HadHeartAttack"
def clustered_plot_counts(ax, df, category_col, heart_col):
    # count (category, heart attack) combinations
    grouped = df.groupby([category_col, heart_col]).size().unstack(fill_value=0)

    # make sure both "No" and "Yes" columns exist and are ordered
    for val in ["No", "Yes"]:
        if val not in grouped.columns:
            grouped[val] = 0
    grouped = grouped[["No", "Yes"]]

    # clustered bar chart with counts and fixed colors
    grouped.plot(kind='bar', ax=ax, color=["tab:blue", "tab:orange"])

    ax.set_title(f"Heart Attack by {category_col}")
    ax.set_ylabel("Number of people")
    ax.tick_params(axis='x', rotation=45)

    ax.legend(["No heart attack", "Heart attack"], title="Heart attack")

fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# 1) Distribution of heart attack (entire population) - counts
counts = df[heart_col].value_counts()

# ensure "No" and "Yes" both present and ordered
for val in ["No", "Yes"]:
    if val not in counts.index:
        counts[val] = 0
counts = counts[["No", "Yes"]]

axes[0, 0].bar(counts.index, counts.values, color=["tab:blue", "tab:orange"])
axes[0, 0].set_title("Heart Attack Distribution (Entire Population)")
axes[0, 0].set_ylabel("Number of people")
axes[0, 0].tick_params(axis='x', rotation=0)

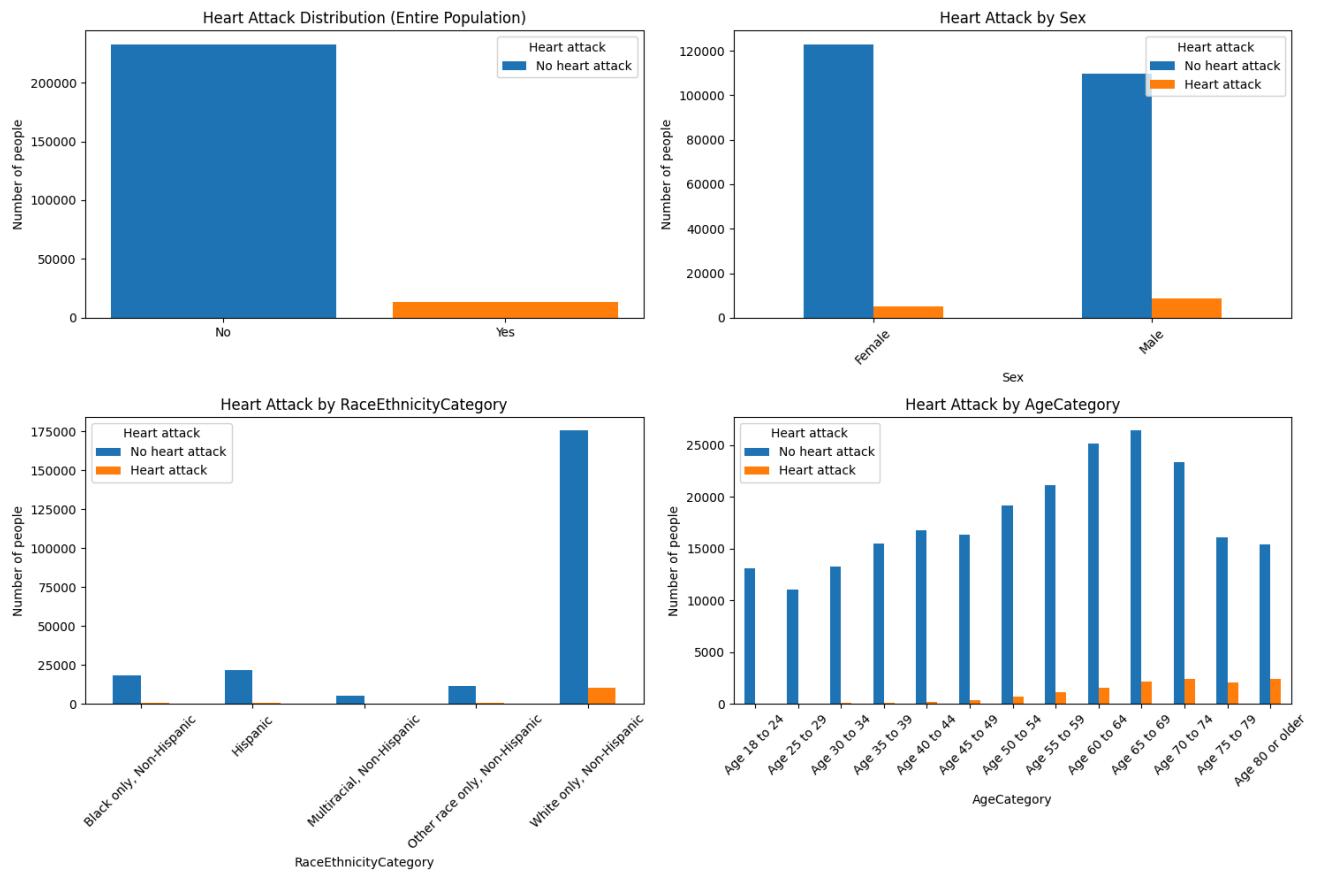
# consistent legend
axes[0, 0].legend(["No heart attack", "Heart attack"], title="Heart attack")

# 2) Heart Attack vs Sex (counts, clustered)
clustered_plot_counts(axes[0, 1], df, "Sex", heart_col)

# 3) Heart Attack vs RaceEthnicityCategory (counts, clustered)
clustered_plot_counts(axes[1, 0], df, "RaceEthnicityCategory", heart_col)

# 4) Heart Attack vs AgeCategory (counts, clustered)
clustered_plot_counts(axes[1, 1], df, "AgeCategory", heart_col)

plt.tight_layout()
plt.show()
```



```

heart_col = "HadHeartAttack"

def plot_factor_vs_heart(ax, df, factor_col, heart_col):
    grouped = df.groupby([factor_col, heart_col]).size().unstack(fill_value=0)

    # make sure columns are ordered ["No", "Yes"] if both exist
    for v in ["No", "Yes"]:
        if v not in grouped.columns:
            grouped[v] = 0
    grouped = grouped[["No", "Yes"]]
    grouped.plot(kind="bar", ax=ax)
    ax.set_title(f"{factor_col} vs {heart_col}")
    ax.set_ylabel("Number of people")
    ax.tick_params(axis="x", rotation=45)
    ax.legend(["No heart attack", "Heart attack"], title="Heart attack")

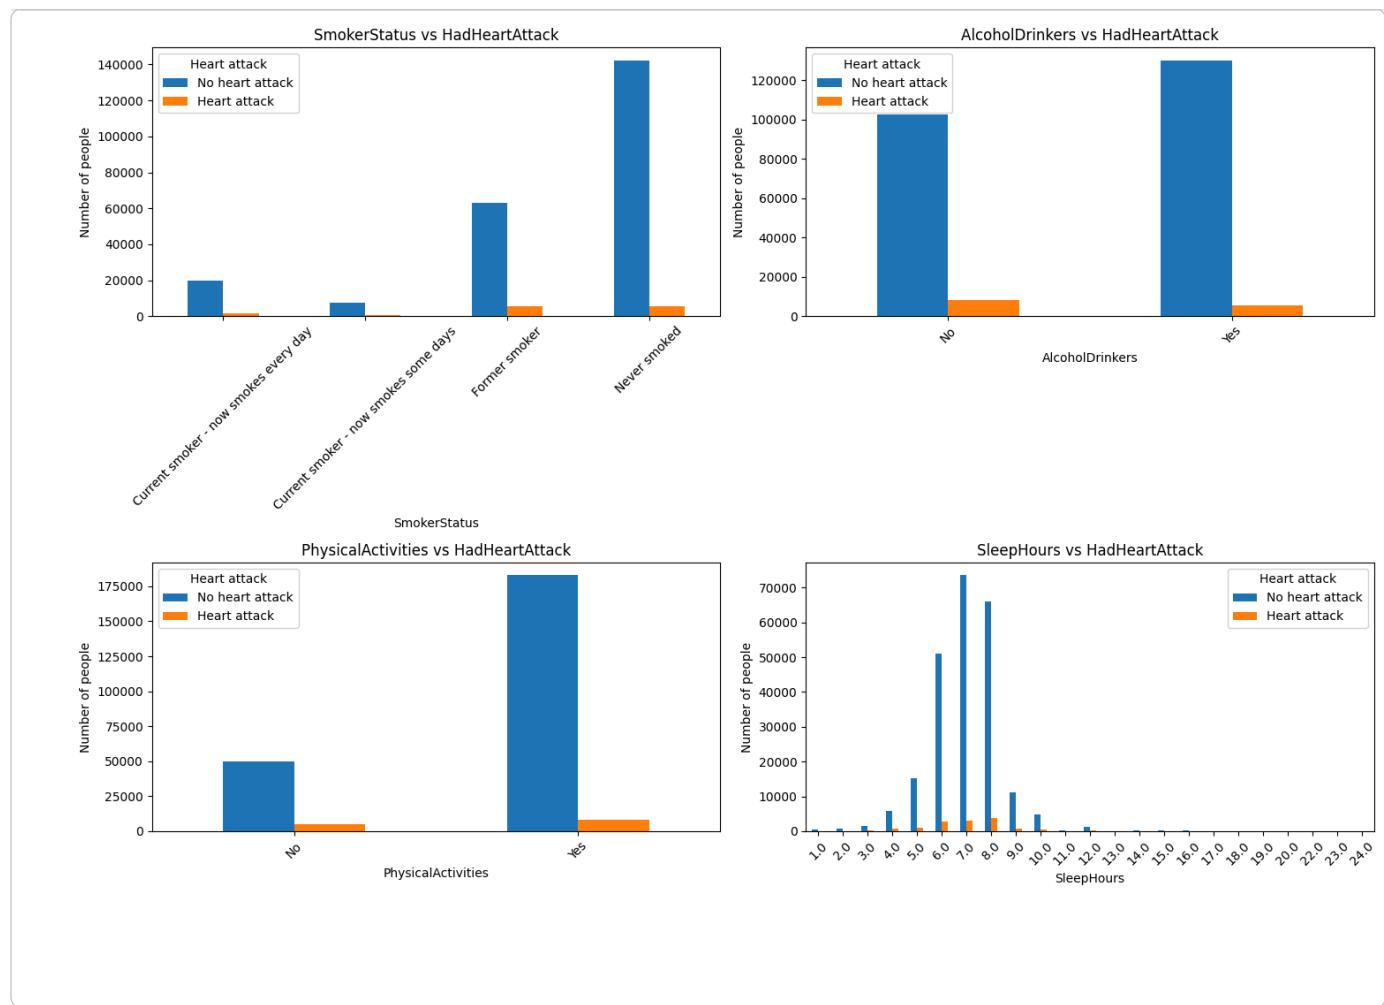
factors = ["SmokerStatus", "AlcoholDrinkers", "PhysicalActivities", "SleepHours"]

fig, axes = plt.subplots(2, 2, figsize=(15, 10))

for ax, factor in zip(axes.flatten(), factors):
    plot_factor_vs_heart(ax, df, factor, heart_col)

plt.tight_layout()
plt.show()

```



```
# Conditions to compare
conditions = ["HadStroke", "HadKidneyDisease", "HadDiabetes", "HadAsthma", "HadSkinCancer", 'HadArthritis',
              'HadDepressiveDisorder', 'DifficultyWalking']

total_counts = []
heartattack_counts = []

for cond in conditions:
    cond_df = df[df[cond] == 'Yes'] # people with the condition
    total = len(cond_df)
    heart_attack = (cond_df['HadHeartAttack'] == 'Yes').sum()

    total_counts.append(total)
    heartattack_counts.append(heart_attack)

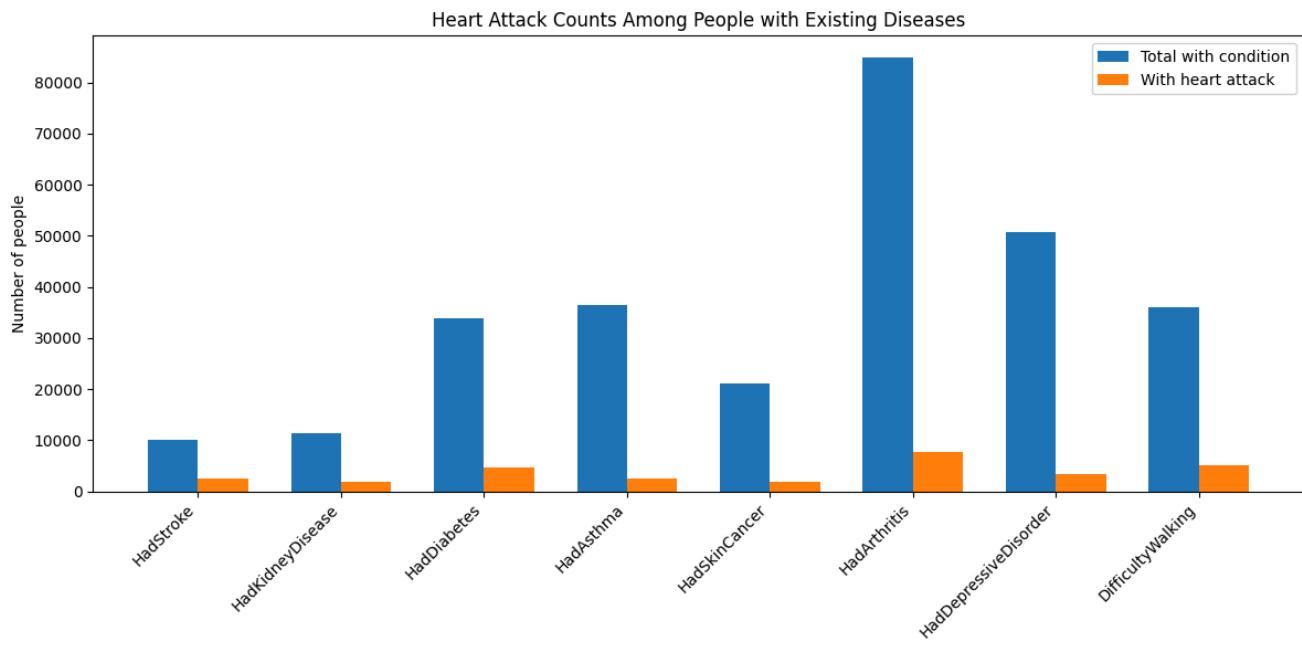
# X locations
x = np.arange(len(conditions))
width = 0.35 # bar width

plt.figure(figsize=(12, 6))

plt.bar(x - width/2, total_counts, width, label='Total with condition')
plt.bar(x + width/2, heartattack_counts, width, label='With heart attack')

plt.xticks(x, conditions, rotation=45, ha='right')
plt.ylabel("Number of people")
plt.title("Heart Attack Counts Among People with Existing Diseases")

plt.legend()
plt.tight_layout()
plt.show()
```



```

import pandas as pd
import matplotlib.pyplot as plt

# Columns for conditions and heart attack
conditions = [
    "HadStroke", "HadKidneyDisease", "HadDiabetes", "HadAsthma",
    "HadSkinCancer", "HadArthritis", "HadDepressiveDisorder",
    "DifficultyWalking"
]

heart_col = "HadHeartAttack"    # adjust if named differently

risk_percentages = []

for cond in conditions:
    # People who have this disease
    cond_yes = df[df[cond] == "Yes"]
    total_with_disease = len(cond_yes)

    if total_with_disease == 0:
        risk = 0
    else:
        heart_attack_count = (cond_yes[heart_col] == "Yes").sum()
        # % heart attack among people who have this disease
        risk = (heart_attack_count / total_with_disease) * 100

    risk_percentages.append(risk)

# ---- Plotting ----

plt.figure(figsize=(12, 6))
plt.bar(conditions, risk_percentages, color="red")

plt.ylabel("Heart Attack Percentage (%)")
plt.title("Risk of Heart Attack Among People With Each Existing Disease")

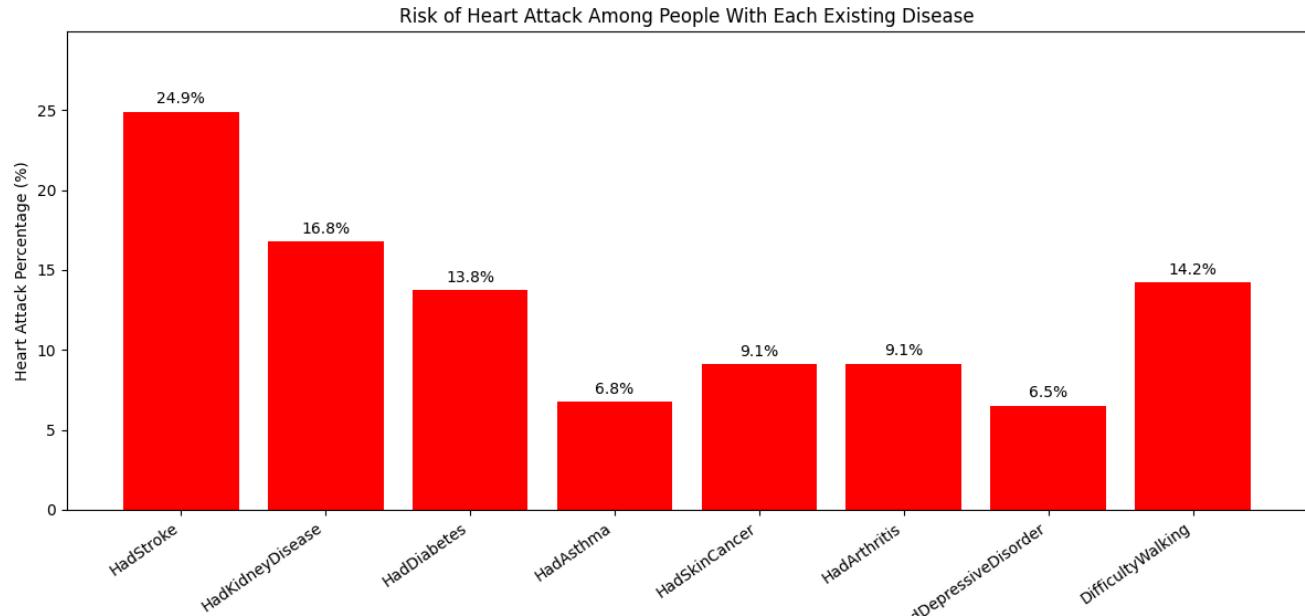
# Fix overlapping labels
plt.xticks(rotation=35, ha='right')

# Add value labels
for i, v in enumerate(risk_percentages):
    plt.text(i, v + 0.5, f"{v:.1f}%", ha='center')

plt.ylim(0, max(risk_percentages) + 5)
plt.tight_layout()

```

```
plt.show()
```



```
# Columns to include in the correlation heatmap
cols = [
    "SmokerStatus",
    "AlcoholDrinkers",
    "PhysicalActivities",
    "SleepHours",
    "HadStroke",
    "HadKidneyDisease",
    "HadDiabetes",
    "HadAsthma",
    "HadSkinCancer",
    "HadArthritis",
    "HadDepressiveDisorder",
    "DifficultyWalking"
]

# Work on a copy with only these columns
corr_df = df[cols].copy()

# Convert non-numeric columns (e.g. Yes/No, categories) to numeric codes
for c in corr_df.columns:
    if not pd.api.types.is_numeric_dtype(corr_df[c]):
        corr_df[c] = corr_df[c].astype("category").cat.codes

# Compute correlation matrix
corr_matrix = corr_df.corr()

# Plot heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f")
plt.title("Correlation Heatmap: Lifestyle Factors and Existing Diseases")
plt.tight_layout()
plt.show()
```



```

## Step 7: Data Encoding
# Map 'Male' to 1 and 'Female' to 0 for 'Sex' column
df['Sex'] = df['Sex'].map({'Male': 1, 'Female': 0})
df['HadHeartAttack'] = df['HadHeartAttack'].map({'Yes': 1, 'No': 0})

# Ordinal encoding for 'AgeCategory' & 'GeneralHealth' (assuming ordering based on age groups & health status)
age_order = ['Age 18 to 24', 'Age 25 to 29', 'Age 30 to 34', 'Age 35 to 39', 'Age 40 to 44', 'Age 45 to 49',
             'Age 50 to 54', 'Age 55 to 59', 'Age 60 to 64', 'Age 65 to 69', 'Age 70 to 74', 'Age 75 to 79',
             'Age 80 or older']
df['AgeCategory'] = df['AgeCategory'].astype('category').cat.set_categories(age_order, ordered=True).cat.codes

health_order = ['Poor', 'Fair', 'Good', 'Very good', 'Excellent']
df['GeneralHealth'] = df['GeneralHealth'].astype('category').cat.set_categories(health_order, ordered=True).cat.codes

# One-Hot Encoding for remaining object columns
object_cols = df.select_dtypes(include='object').columns
df = pd.get_dummies(df, columns=object_cols, drop_first=True)
# object_cols list of column names that are categorical, usually with dtype object
# get_dummies turns each categorical col into multiple 0/1 columns (one-hot encoding--eg., Gender_Female Gender_Male)
# Boolean columns (T/F) have dtype = bool, so they are not in object_cols & get_dummies simply never touches them

```

df

```
Sex GeneralHealth PhysicalHealthDays MentalHealthDays SleepHours HadHeartAttack AgeCategory HeightInMetre
```

342	0	3	4.0	0.0	9.0	0	9	1.
343	1	3	0.0	0.0	6.0	0	10	1.
345	1	3	0.0	0.0	8.0	0	11	1.
346	0	1	5.0	0.0	9.0	0	12	1.
347	0	2	3.0	15.0	5.0	0	12	1.
...
445117	1	3	0.0	0.0	6.0	0	8	1.
445123	0	1	0.0	7.0	7.0	0	1	1.
445124	1	2	0.0	15.0	7.0	0	9	1.
445128	0	4	2.0	2.0	7.0	0	6	1.
445130	1	3	0.0	0.0	5.0	1	10	1.

246013 rows × 108 columns

```
# Select boolean columns & convert to integer (True to 1, False to 0)
boolean_cols = df.select_dtypes(include='bool').columns
for col in boolean_cols:
    df[col] = df[col].astype(int)
```

```
## Step-8: Splitting Dataset- Declare feature vector and target variable
X = df.drop('HadHeartAttack', axis=1)
y = df['HadHeartAttack']

# Split data into separate training and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step-9: Training RF Model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
```

▼ RandomForestClassifier ⓘ ?

```
RandomForestClassifier(random_state=42)
```

```
# Evaluation Metrics
y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy*100:.2f}%"") #accuracy is high because the majority class dominates, not because the mode
print(f"Precision: {precision*100:.2f}%"") # Precision: 62.36%
print(f"Recall: {recall*100:.2f}%"") # Recall: 12.78%, Out of all REAL heart attack cases, our model catches only 12%.
# It shows the model:
# Is biased toward predicting No heart attack
# Is failing to learn the minority class
# Suffers heavily from class imbalance
# Our model catches only 12.78% of actual heart attack cases.
print(f"F1 Score: {f1*100:.2f}%"")
```

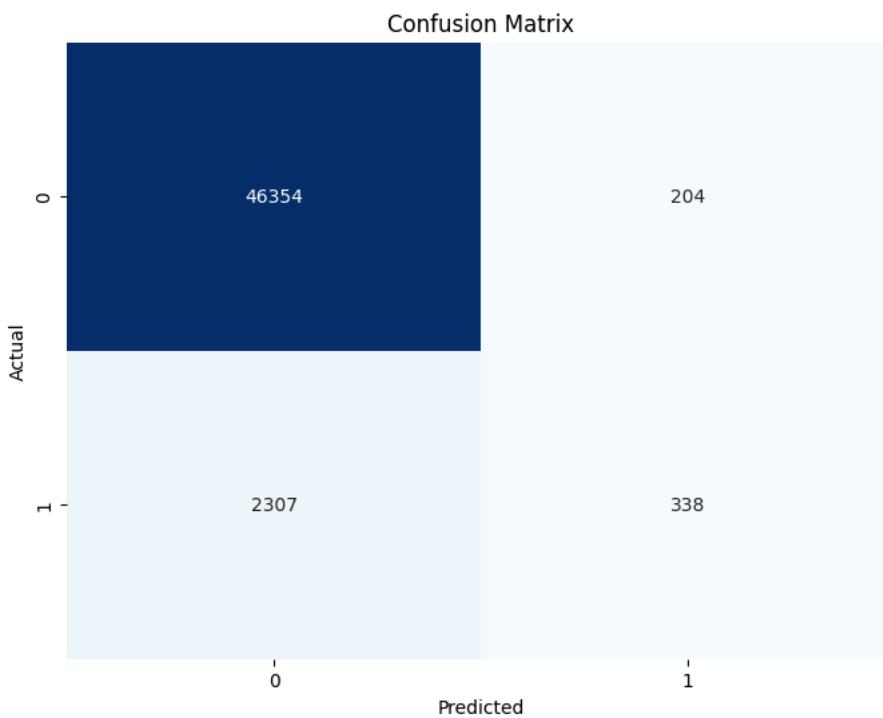
Accuracy: 94.90%
Precision: 62.36%
Recall: 12.78%
F1 Score: 21.21%

```
class_report = classification_report(y_test, y_pred)
print("Classification Report:\n", class_report)
```

Classification Report:		precision	recall	f1-score	support
0	0.95	1.00	0.97	46558	
1	0.62	0.13	0.21	2645	
accuracy				0.95	49203
macro avg	0.79	0.56	0.59	49203	
weighted avg	0.93	0.95	0.93	49203	

```
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
from imblearn.over_sampling import SMOTE
## Splitting Dataset and SMOTE
X = df.drop('HadHeartAttack', axis=1)
y = df['HadHeartAttack']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Apply SMOTE to the training data
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Apply SMOTE to the testing data
X_test_smote, y_test_smote = smote.fit_resample(X_test, y_test)
##Training RF Model SMOTE
model_smote = RandomForestClassifier(random_state=42)
model_smote.fit(X_train_smote, y_train_smote)

## Evaluation Metrics using SMOTE Data
y_pred_smote = model_smote.predict(X_test_smote)

accuracy_smote = accuracy_score(y_test_smote, y_pred_smote)
precision_smote = precision_score(y_test_smote, y_pred_smote)
recall_smote = recall_score(y_test_smote, y_pred_smote)
f1_smote = f1_score(y_test_smote, y_pred_smote)
```

```
print("\nEvaluation with SMOTE:")
print(f"Accuracy: {accuracy_smote*100:.2f}%")
print(f"Precision: {precision_smote*100:.2f}%")
print(f"Recall: {recall_smote*100:.2f}%)
```

```
Evaluation with SMOTE:
Accuracy: 91.06%
Precision: 0.7500
Recall: 0.7500
```